

An Empirical Study of Double Descent in Polynomial Regression

Konark

ITSM EKONARK@GMAIL.COM

School of Computer Science
Constructor University Bremen
28759 Bremen, Germany

Editor: Academic Skills in Computer Science

Abstract

Machine-learning textbooks teach that as a model gets more complex, its error on new data first goes down and then back up. The result is a U-shape. This is not the whole story. If you keep adding parameters past the point where the model fits the training data exactly, the error often goes *down again*. This is called *double descent*. I show it with about as simple an experiment as one can ask for: fitting polynomials to a wiggly curve. With plenty of training data, I get the textbook U-shape. With little training data, a sharp spike appears at the point where the polynomial has just enough wiggles to hit every training point—and then the error drops a second time as the polynomial keeps growing. Splitting the spike into bias and variance shows that both blow up at the same time. Adding even a tiny amount of ridge regularisation makes the spike disappear. **Keywords:** double descent, polynomial regression, bias–variance, overparameterisation, ridge regression

1 Introduction

Almost every machine-learning class starts the same way. A model that is too simple cannot capture the truth—it has high *bias*. A model that is too complex chases random noise—it has high *variance*. Plot test error against model size and you get a U: bad on the left, bad on the right, best somewhere in the middle. This has been the standard story for over thirty years (Geman et al., 1992; Hastie et al., 2009).

The story does not match what people actually do today. Modern neural networks have millions of parameters and only a few thousand training examples. They fit the training data perfectly. The U-shape predicts they should be useless on new data, but in practice they work very well.

In 2019, Belkin et al. (2019) gave a name to what is really going on: *double descent*. The error follows the U-shape for small models. Then it spikes—sharply—at the point where the model has just enough parameters to fit the training set exactly. Then it goes down again as the model keeps growing. Researchers have since seen the same shape in many other kinds of models, including deep neural networks (Nakkiran et al., 2021a).

This paper has one small goal: to show double descent in the simplest setting I could find. No neural networks, no heavy theory. Just fitting polynomials to a one-dimensional toy problem. The experiments make three points. **(i)** The choice of basis matters. Plain monomials $1, x, x^2, \dots$ blow up numerically for moderate model sizes, so I use a more stable family. **(ii)** The spike at the threshold is caused by *both* bias and variance blowing up, not by variance alone as is sometimes claimed. **(iii)** Even a tiny ridge penalty removes the spike, so double descent is fragile. All the code, figures, and $\text{\LaTeX}$ sources are released alongside this paper so anyone can rerun the experiments.

2 Related Work

The classical bias–variance trade-off is described in any standard textbook (Hastie et al., 2009) and was given a famous treatment in the neural-network setting by Geman et al. (1992). Belkin et al. (2019) brought double descent to the attention of the machine-learning community, observing it in random Fourier features, decision trees, and small neural networks. Nakkiran et al. (2021a) demonstrated it cleanly in modern deep networks. For linear models, Hastie et al. (2022) now provide a fairly complete asymptotic theory, and Nakkiran et al. (2021b) prove that with the right amount of ridge regularisation the spike disappears entirely. I do not engage with this theory mathematically. I just re-derive its qualitative predictions in a small experiment that anyone can run on a laptop.

3 Experimental Setup

The setup is intentionally simple. There are only three pieces: data, a model, and a way to fit the model.

Data. I draw n random points between 0 and 1. The label of each point is a smooth wave plus a bit of Gaussian noise:

$$y = \sin(2\pi x) + 0.3x + \varepsilon, \quad \varepsilon \sim \mathcal{N}(0, 0.3^2). \quad (1)$$

The test set is 2000 fresh points drawn from the same recipe.

Model. Polynomial regression with p features. The natural choice would be $1, x, x^2, \dots, x^{p-1}$, but those columns become almost identical for large p and the math blows up numerically. I use *Legendre polynomials* instead. These are just a different family of polynomials that stay numerically well-behaved as p grows. NumPy provides them out of the box (`numpy.polynomial.legendre.legvander`).

Fit. Standard least squares. When the model has more parameters than data points ($p > n$), there are infinitely many fits that pass through every training point. I pick the one with the smallest coefficients—this is what `numpy.linalg.lstsq` returns. Some experiments also add a small *ridge* penalty λ that discourages large coefficients.

Reporting. I run each experiment many times (30–400) with fresh random data and report averages or medians. All random seeds are fixed in the code, so every figure is bit-for-bit reproducible.

4 Experiments

4.1 The classical regime: a tame U-curve

I first check that polynomial regression in the classical regime gives the textbook U. With $n = 200$ training points and polynomial degrees up to 24 (so $p \ll n$), the training error decreases monotonically, as it has to (Figure 1). The test error reaches a minimum at degree ≈ 5 , which roughly matches the smoothness of the target (1), and slowly drifts up as the model starts chasing noise. There are no surprises here. This is the trade-off described by Geman et al. (1992).

4.2 Double descent across the interpolation threshold

I now shrink the training set to just $n = 30$ points and let the number of features p run from 1 to 200. With so little training data, p now sweeps through—and well past—the *interpolation threshold* $p = n$, where the polynomial has just enough parameters to fit every training point exactly. Figure 2 shows what happens.

The shape is clear. For $p < n$, the test error follows the familiar U-curve and reaches a small minimum near $p \approx 6$. As p approaches the threshold the test error climbs sharply and reaches a peak more than *ten orders of magnitude* above the noise floor at $p = n$. Past the threshold the training error drops to numerical




Figure 1: Classical U-shaped trade-off. With $n = 200$ training points and polynomial degrees up to 24 (so $p \ll n$), the test error is minimised at degree ≈ 5 and slowly creeps up after that. Shading: ± 1 standard deviation over 30 independent training samples.

zero, and the test error *decreases* towards a level only a little worse than the first descent’s minimum. This is the canonical double-descent shape predicted by recent theory (Hastie et al., 2022) and observed in many other models (Nakkiran et al., 2021a).

It is worth pausing on *why* the second descent happens at all. With $p > n$, infinitely many polynomials fit the data exactly. The algorithm picks the one with the smallest coefficients, which is the smoothest available choice. As p grows, the set of admissible fits gets bigger, and the smoothest one becomes smoother in the directions that matter, so the test error goes down. Loosely: more parameters give the algorithm *more room* to find a smooth fit.

4.3 Bias and variance both blow up at the spike

Where does the spike at $p = n$ come from? A common informal explanation is that variance explodes. Some authors instead point at bias. To settle the question empirically, I split the test error into squared bias and variance using 400 independent training samples (Figure 3).

The decomposition shows that *both* terms peak around $p = n$. Squared bias is large near the threshold because the average prediction (over training samples) is itself an erratic curve between the training points, far from the smooth target. Variance is large because tiny changes to the training set lead to wildly different fits. Both effects come from the same underlying cause: the design matrix becomes nearly singular at $p = n$. Once p grows beyond n , the design matrix recovers and both terms decay.

4.4 Ridge regularisation removes the spike

The minimum-coefficient solution above corresponds to ridge regression with $\lambda = 0$. Nakkiran et al. (2021b) prove that with the right amount of ridge, the test error is monotone decreasing in p and the spike disappears entirely. Figure 4 reproduces this finding.

Even a tiny ridge penalty, $\lambda = 10^{-4}$, brings the peak down from about 10^{18} to about 10^1 . With $\lambda = 10^{-2}$ the curve is essentially flat across the threshold. With $\lambda = 1$ the test error is monotone in p . This matches the theoretical prediction of Nakkiran et al. (2021b) and gives a useful practical takeaway: in linear regression, double descent only appears when there is no explicit regularisation at all.




Figure 2: Double descent in least-squares polynomial regression. With $n = 30$ training points, the test error follows the U-shape for $p < n$, peaks sharply at the threshold $p = n = 30$, and *descends a second time* as the model becomes overparameterised. Solid lines: medians over 80 training samples; shading: interquartile range.

5 Discussion

What the experiments do and do not show. I have demonstrated double descent in the cleanest setting I could put together, but the experiments are simple by design. Two limitations are worth flagging. First, the cleanliness of the curve depends on a well-behaved basis. The same experiment with raw monomials produces an unreadable plot dominated by floating-point error. Second, the setup is one-dimensional, and the smooth wiggly target in (1) can be represented exactly by a sufficiently large polynomial, so the bias term in the overparameterised regime decays especially fast. In higher dimensions, or when the target function does not lie inside the model class, the second descent is less aggressive (Hastie et al., 2022).

Why this matters in practice. Two takeaways are worth keeping in mind. First, the intuition that “more parameters than data points” is automatically pathological is wrong; the minimum-coefficient solution implicitly regularises and can generalise just fine. Second, if you suspect you are operating near the threshold, even a small amount of explicit regularisation removes the risk of catastrophic test error.

Connection to deep learning. The phenomenon shown here in linear regression is closely related, though not identical, to “model-wise” double descent in deep neural networks (Nakkiran et al., 2021a). In both cases the peak occurs where the model has just enough capacity to fit the training labels. In both cases regularisation, early stopping, or more data tends to mitigate it. The neural-network case has additional sources of implicit regularisation coming from the optimiser (e.g. stochastic gradient descent), which the linear-regression analogue does not capture.

6 Conclusion

I have given a small, self-contained empirical demonstration of double descent in least-squares polynomial regression. Three messages emerge. A well-behaved basis is needed to see the phenomenon cleanly. The spike at the interpolation threshold reflects an explosion in *both* bias and variance, not just one of them. And




Figure 3: Monte Carlo bias–variance decomposition with $n = 30$ training points and 400 independent training samples. *Both* squared bias and variance peak around the threshold $p = n = 30$; neither alone explains the spike.

even a small ridge penalty makes the spike disappear. I hope this short paper works as a useful pedagogical companion to the more sophisticated theoretical treatments cited above.

Reproducibility

The full source code, the LaTeX sources of this paper, and all the figures are released as a single archive accompanying this submission. Each experiment runs in less than a minute on a laptop. Random seeds are fixed in the script.

Acknowledgments and Disclosure of Funding

The author thanks the instructors of the Academic Skills in Computer Science course at Constructor University Bremen for the opportunity to write this paper. No external funding supported this work, and the author declares no competing interests.

References

- Mikhail Belkin, Daniel Hsu, Siyuan Ma, and Soumik Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32):15849–15854, 2019.
- Stuart Geman, Elie Bienenstock, and René Doursat. Neural networks and the bias/variance dilemma. *Neural Computation*, 4(1): 1–58, 1992.
- Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning*. Springer, 2 edition, 2009.
- Trevor Hastie, Andrea Montanari, Saharon Rosset, and Ryan J. Tibshirani. Surprises in high-dimensional ridgeless least squares interpolation. *The Annals of Statistics*, 50(2):949–986, 2022.
- Preetum Nakkiran, Gal Kaplun, Yamini Bansal, Tristan Yang, Boaz Barak, and Ilya Sutskever. Deep double descent: Where bigger models and more data hurt. *Journal of Statistical Mechanics: Theory and Experiment*, 2021(12):124003, 2021a.
- Preetum Nakkiran, Prayaag Venkat, Sham Kakade, and Tengyu Ma. Optimal regularization can mitigate double descent. In *International Conference on Learning Representations (ICLR)*, 2021b.




Figure 4: Effect of ridge regularisation on double descent. Even $\lambda = 10^{-2}$ flattens the spike at $p = n$ almost completely; $\lambda = 1$ gives a curve that is essentially monotone in p .